


Automated Detection and Morphological Classification of Mast Cells in Bone Marrow Aspirates Using Deep Learning



Bachelor Thesis

by
Florian Vollmer

Bachelors 2023

Submission date: 2026-07-02

Study Direction: Applied Digital Life Science

Supervisor:

Dr. Stefan Glüge

ZHAW Life Sciences und Facility Management, Wädenswil

Cross Corrector:

Prof. Dr. Robert Vorburger

ZHAW Life Sciences und Facility Management, Wädenswil

Imprint

Recommended Citation:

Florian Vollmer (2026). *Automated Detection and Morphological Classification of Mast Cells in Bone Marrow Aspirates Using Deep Learning*

.

Keywords: Deep Learning, Cell Detection, Evaluation Metrics, Digital Lab and Production, Clinical Image Processing

Abstract

English Version

German Version

Acknowledgement

I would like to express my sincere gratitude to Dr. Stefan Glüge for setting up the project for this thesis and for his valuable support and guidance throughout the writing and compilation process of this thesis. My appreciation also goes to Robin Kosche for showing me around the USZ and introducing me to the subject.

Table of Contents

List of Abbreviations	6
1. Introduction	7
2. Methodology	8
2.1. Biology of Mast Cells and Systemic Mastocytosis	8
2.2. Technical	9
3. Results	25
3.1. Model Results and Statistical Evaluation	25
3.2. Quantitative and Qualitative Evaluation	29
3.3. Evaluation Visualization	33
3.4. Real World Usage	37
4. Discussion	39
4.1. Interpretation of Results	39
4.2. Clinical Relevance	42
4.3. Error Analysis	43
5. Conclusion	44
5.1. Summary of Key Findings	44
5.2. Limitations and Challenges	44
5.3. Future Research Directions	44
6. Declaration on AI Assistance	45
7. References	46
8. Lists	47
List of Figures	47
List of Tables	48

9. Attachments	49
----------------------	----

List of Abbreviations

- **AI:** Artificial Intelligence
- **AP:** Average Precision
- **CNN:** Convolutional Neural Network
- **COCO:** Common Objects in Context
- **CV:** Cross Validation
- **CVAT:** Computer Vision Annotation Tool
- **DoE:** Design of Experiment
- **FN:** False Negative
- **FP:** False Positive
- **GT:** Ground Truth
- **IoU:** Intersection over Union
- **LOPO:** Leave-One-Patient-Out
- **MC:** Mast Cells
- **mAP:** Mean Average Precision
- **NMS:** Non-Maximum Suppression
- **RF-DETR:** Roboflow Detection Transformer
- **SM:** Systemic Mastocytosis
- **TN:** True Negative
- **TP:** True Positive
- **USZ:** University Hospital Zurich (Universitätsspital Zürich)
- **WHO:** World Health Organization
- **YOLO:** You Only Look Once

1. Introduction

2. Methodology

2.1. Biology of Mast Cells and Systemic Mastocytosis

2.2. Technical

General Data Information

The dataset for this study consists of digitized bone marrow aspirate smear images provided by the University Hospital Zurich (USZ). All samples were collected from patients with confirmed or suspected systemic mastocytosis. Image tiles were extracted from whole-slide scans at a fixed resolution of 512x512 pixels, each displaying a sub-section of the whole smear. Two morphological classes were labeled for this study: Atypisch (atypical mast cells) and Normal (normal mast cells), corresponding directly to the cell types relevant to the WHO minor diagnostic criterion for SM.

The corpus grew across four acquisition batches spanning 53 patients. The initial dataset covered a single patient (P1) and served as the pilot for the annotation and training pipeline. A second batch extended the corpus to 15 patients (P1-P15), contributing 1,564 annotated images with 1,288 Atypisch and 391 Normal instances, reflecting a raw class ratio of approximately 3.3:1. A third batch added patients P16 through P53, bringing the full corpus to 53 patients (1,857 annotated images, 1,476 Atypisch, 511 Normal, raw ratio 2.9:1). Because the training data derives almost entirely from confirmed SM patients, Atypisch is the dominant class throughout. At the P1-P8 scale the raw ratio was 14.4:1; by P1-P15 this had dropped to 3.3:1 due to the inclusion of patients P6, P8, and P9, which are rich in Normal annotations.

Data Acquisition and Annotation

The annotation pipeline was carried out by Robin Kosche at USZ using CVAT (Computer Vision Annotation Tool). Rather than annotating all 53 patients manually from scratch, an iterative model-assisted annotation strategy was used to make the labeling process tractable at the scale of tens of thousands of image tiles.

The first patient (P1) was annotated entirely by hand. Bounding boxes were drawn around each visible mast cell. This initial patient had only Atypisch cells. This dataset was used to train the initial domain-specific detection model, DL_Modell_FV.pt, a YOLO11n checkpoint trained on P1 data.

For patients P2 through P15, CVAT's model inference integration was used with DL_Modell_FV.pt as the backend. The model generated candidate bounding boxes for each image tile, which the annotator then reviewed, accepted, corrected, or rejected. Starting from this batch, Normal cells were also annotated. Tiles where the model produced detections on regions that contained no mast cell were recorded separately in per-patient Excel files. These confirmed false positives were later incorporated into training as a curated set of hard-negative examples.

For patients P16 through P53, the same workflow was repeated using a second-generation model (DL_Modell_FV_v2.pt), trained on the P1-P15 corpus, as the inference backend. It is worth noting that a subset of patients in this batch contained no annotatable mast cells at all: the model produced candidate detections, but all were rejected by the annotator as false positives. These patients contributed no positive labels to the corpus but supplied a large number of confirmed hard negatives. The annotations from the full P1-P53 corpus were then used to train a third-generation model, DL_Modell_FV_v3.pt, completing the bootstrapped annotation cycle at the current project scope.

This iterative approach, where each annotation round produces a labeled dataset used to train a better model for the next round, allowed the corpus to grow from one patient to 53 without

requiring fully manual annotation of all 53 cases. The confirmed false positives collected at each round provided a particularly valuable training signal: these are the specific tiles that the model of that generation found difficult to reject, making them targeted hard negatives for the next training cycle.

To our knowledge, this bootstrapped annotation strategy is not documented for mast cell detection in the existing literature. The systematic review by Ghete et al. [1], covering AI-based bone marrow cell classification from 2019 to 2024, found that the majority of published studies applied their models to pre-existing datasets rather than developing annotation pipelines alongside the model. In this project, annotation and model development were tightly coupled by design: the clinical partner’s annotation capacity was directly extended by each generation of model, and the confirmed false positives produced during annotation became training signal for the next generation. This coupling is what made a 53-patient corpus feasible within the scope of a single thesis.

The final corpus assembled across all three annotation rounds is summarized in Table 1.

Table 1: Full corpus summary across P1-P53.

Category	Count
Total annotated images	1’857
Atypisch cell labels	1’476
Normal cell labels	511
Raw class ratio (Atyp:Norm)	2.9 : 1
Confirmed FP/Background images	4’940

Data Preprocessing

Two distinct preprocessing approaches were applied across the experimental timeline, with a significant revision midway through the project.

Pre-applied augmentation (Phase 1 and early Phase 2). For the initial experiments, a preprocessing pipeline was implemented using the `albuminations` library. Each annotated image was transformed once before training with random flips, brightness and contrast adjustments, and affine rotations of up to 15 degrees. The augmented copies were saved to disk and used as fixed training input. This approach proved ineffective: because transforms were applied once and stored, every training epoch saw the same augmented copies. The augmented images extended the dataset by a fixed factor but provided no additional regularization benefit after the first epoch, which is the only epoch where the transforms are novel to the model. This limitation was identified during Phase 2 and motivated the switch to on-the-fly augmentation.

On-the-fly augmentation (Phase 3 onward). Starting from Phase 3, augmentation was moved into the YOLO training loop by enabling `augment=True` in the `model.train()` call. This applies a fresh random transform to each image at every epoch, meaning oversampled copies in the training file list present different pixel content to the model across epochs. Two augmentation parameters required explicit configuration for this domain. The mosaic augmentation was disabled (`mosaic=0.0`): YOLO’s mosaic mode composites four images into a single training sample, which is appropriate for multi-object scene datasets like COCO but is counterproductive for 512x512

bone marrow tiles, where the spatial relationship between a candidate cell and its immediate local context is part of the morphological signal the model needs to learn. Vertical flip probability was set to 0.5 (the YOLO default is 0), consistent with the biological reality that bone marrow mast cells have no canonical vertical orientation and that the model should not learn a vertical bias from the training data.

The rotation augmentation parameter (degrees) was treated as a tunable hyperparameter during Phase 3. A systematic tinkering matrix identified degrees=5 as optimal at the Phase 3 corpus scale, providing approximately 4-5 percentage points of improvement in Normal-class recall over degrees=0. The effect is biologically justified: mast cells present in any orientation in aspirate smears, and moderate rotation augmentation regularizes the model without distorting morphological cues that distinguish Atypisch from Normal.

Model Selection and Architecture

YOLO11 (Ultralytics) was selected as the object detection framework. The architecture performs detection and classification in a single forward pass, making it computationally efficient for integration into clinical workflows where inference throughput matters. The Ultralytics Python API provides a consistent interface for training, fine-tuning, evaluation, and inference, and supports checkpoint initialization from custom pretrained weights, which was central to the transfer learning strategy described below. Comparable studies applying YOLO-family models to cell detection in hematological microscopy report mAP50 values in the range of 77-94% on well-represented cell types [2] [3], placing the architecture in the performance range expected for this domain.

Among the available YOLO11 model sizes, the nano variant (yolo11n, 2.6M parameters) was selected. The selection was validated experimentally: a test using the large variant (yolo11l, 43M parameters) on the P1-P8 corpus under identical training conditions produced $R_{\text{Normal}}=0.695$ compared to $R_{\text{Normal}}=0.840$ for yolo11n. The larger model's classification head, which is randomly re-initialized when training with a different number of output classes than the pretrained checkpoint, has proportionally more free parameters to adapt from scratch. With only 170-220 training images per LOPO (Leave-One-Patient-Out) fold, these parameters tend to overfit to the majority class rather than learn generalizable features for the rare Normal class. The corpus size at this stage was insufficient for a 43M-parameter model; yolo11n was confirmed as the appropriate architecture for the available data.

Throughout the project, all training runs started from the generic COCO-pretrained yolo11n.pt checkpoint. The single exception was a set of experiments during Phase 2 in which DL_Modell_FV.pt, a YOLO11n model previously trained on the P1 single-patient dataset (nc=1, one class), was used as the starting point for continued training on the two-class P2 dataset. The reasoning was that domain-specific pretraining might give the model a better starting point for bone marrow aspirate images compared to COCO features. In practice, this approach introduced a structural complication: because DL_Modell_FV.pt was trained with a single output class and the P2 training required two classes (Atypisch and Normal), the Detect head's classification branch (cv3, the final 1x1 convolution layer) had to be discarded and randomly re-initialized. The bounding-box regression branch (cv2) transferred cleanly, but the randomly initialized cv3 layer produced noisy classification gradients in early training epochs. At the default learning rate of $lr=0.01$, these gradients were large enough to partially overwrite the pretrained neck features, causing systematic underperformance rather than the expected improvement. The full account

of this failure and its resolution is described in the training strategy section. After Phase 2, the project returned to yolo11n.pt as the standard base for all subsequent training.

Model Training and Finetuning Strategy

The training strategy developed across five phases, each responding to specific failures and insights from the previous phase. The following account describes the major decision points rather than every individual run; a complete run table is provided in the Results section. The phase structure corresponds to successive expansions of the training corpus and methodology: Phase 1 covered a single pilot patient, Phase 2 moved to a two-class annotated dataset and explored data composition, Phase 3 introduced multi-patient LOPO cross-validation and on-the-fly augmentation on an 8-patient corpus, Phase 4 expanded to 15 patients, and Phase 5 targets the full 53-patient corpus. Individual experiment runs within each phase are identified by a label combining the phase number and a letter (e.g., P2-C, P3-A), which matches the experiment log maintained throughout the project.

Phase 1: Single-patient proof of concept and hyperparameter search.

Phase 1 covered a sequence of experiments on the P1 dataset (428 images, single patient, Atypisch class only) using 5-fold stratified cross-validation and the pre-applied augmentations augmentation pipeline described above. The dataset consisted of annotated positive tiles alongside hand-curated negative tiles (confirmed empty regions), all at 512x512 pixels.

The initial runs used yolo11n.pt (COCO pretrained). To test whether a larger model would learn better features for this domain, yolo11l (43M parameters) was then substituted and run in three separate configurations, including extended epoch budgets (200 epochs) and added augmentation. In each case, the large model did not outperform yolo11n on this single-patient dataset, and the project returned to yolo11n for all further Phase 1 work. This result prefigured P3-E; the mechanism is described in §Model Selection and Architecture.

With yolo11n established as the appropriate architecture, hyperparameter tuning was run using YOLO's built-in `model.tune()` function across three increasingly intensive rounds: 30 iterations x 50 epochs, 300 iterations x 50 epochs, and finally 300 iterations x 100 epochs. The output of the final tuning round was a `best_hyperparameters.yaml` file with optimized values including `lr0=0.0062`, `momentum=0.791`, `mosaic=0.916`, and `degrees=0.0`. A training run applying these tuned hyperparameters completed Phase 1. The per-fold results are reported in the Results section.

The `best_hyperparameters.yaml` produced by the Phase 1 tuner is significant beyond Phase 1: it was later applied to fine-tuning on P2 data (P2-E) and caused catastrophic forgetting, collapsing performance across all five folds. The failure was later attributed to a recipe mismatch: the tuner had optimised on a from-scratch yolo11n training run, and its output (particularly the high mosaic value and `degrees=0.0`) was calibrated for scratch training from COCO, not for fine-tuning in this clinical domain. The Phase 1 yaml was permanently retired from the fine-tuning pipeline as a result.

Following the 5-fold training, an inference run was conducted on the full first-patient slide image set, approximately 20,000 tiles, excluding all training images by filename. The resulting confidence distribution provided an early look at the model's behavior at deployment scale and informed the direction of subsequent annotation and training work.

Phase 2: The domain-specific checkpoint experiment and the learning rate problem.

Phase 2 began with two runs using the standard yolo11n.pt base on the annotated P2 dataset (P2-A, P2-B), but both were corrupted by a concurrent SLURM job collision described below. The third run (P2-C) introduced DL_Modell_FV.pt as the starting point, testing whether the P1-trained domain-specific model would give the detector a better foundation than COCO-pretrained weights. The result was a characteristic failure pattern: high precision combined with very low recall, the signature of a model that makes few predictions and misses most cells. The cause was the interaction between the randomly re-initialized cv3 head and the default learning rate of $lr=0.01$: the noisy gradients from the freshly initialized classification branch were large enough to partially overwrite the pretrained neck features during early epochs, leaving the model worse than yolo11n.pt would have been from scratch.

Reducing the learning rate to $lr=0.001$ recovered the performance lost to this overwriting effect. Freezing the first ten backbone layers ($freeze=10$) was introduced alongside the learning rate change to prevent the small dataset from further disrupting features that COCO pretraining had established. This combination ($lr=0.001$, $freeze=10$) defined the reference configuration for the remainder of Phase 2. However, the conclusion drawn from the full Phase 2 experiment set was that using DL_Modell_FV.pt as a starting point added complexity without a measurable benefit over yolo11n.pt, and all subsequent phases returned to the standard COCO-pretrained base.

Before expanding to the full P1-P15 corpus that would produce the second model version, a 3×4 factorial design-of-experiment (DoE) grid was run to identify the best data composition strategy at Phase 2 scale. The goal was practical: the hyperparameter and data configuration choices made here would carry forward into the P1-P15 training, and a better second model version meant fewer manual corrections for the annotation partner at USZ when working through patients P16-P53. The grid varied confirmed FP negative oversampling (1x, 2x, 3x) and random background tile ratio (0, 1, 2, 3) across twelve configurations. The mAP50 spread across all twelve DoE configurations was smaller than the cross-fold standard deviation within any single run. Data composition manipulations at single-patient scale could not produce statistically interpretable performance differences; the limiting factor was data scarcity rather than class balance. Background tiles uniformly hurt performance at every FP oversampling level, and this finding was consistent enough to permanently set the random background ratio to zero in all subsequent runs. The DoE finding reframed the problem: improving recall required more patients and better augmentation, not a different ratio of existing data categories. The transferable lesson for rare-event medical imaging is discussed in the Discussion.

Phase 2 also exposed an infrastructure failure specific to the HPC environment. Submitting two SLURM jobs concurrently caused both jobs to share the same temporary workspace directories, because directory paths were derived from fold indices without job-level namespacing. Concurrent writes to shared directories corrupted three of five folds in each job, detectable by fold results that were bit-for-bit identical between the two affected runs. The fix applied \$SLURM_JOB_ID namespacing to all temporary directories and derived the output project directory name from the job name set at sbatch submission time, making concurrent jobs structurally independent.

Phase 3: Multi-patient LOPO and on-the-fly augmentation.

Phase 3 restructured the training and evaluation design around three simultaneous changes: expansion to 8 patients, adoption of Leave-One-Patient-Out cross-validation, and the switch to on-the-fly augmentation.

LOPO was adopted to eliminate the patient-level leakage present in stratified k-fold splits: holding out all images from one patient per fold makes each fold an evaluation of genuine cross-patient generalization. The full rationale and its relation to external validation standards are in Section .

The combined effect of these three changes was substantial. The Phase 3 baseline run (P3-A: 8 patients, yolo11n.pt, augment=True, mosaic=0.0, cls=1.0) achieved the strongest results seen to that point in the project, with cross-fold standard deviation markedly lower than the Phase 2 DoE despite the harder evaluation protocol: the more diverse training corpus more than compensated for the stricter holdout design. Full results are reported in the Results section.

With the same goal of finding the best possible training recipe before the P1-P15 training that would produce the second model version, a tinkering matrix of 14 parameter combinations explored rotation degrees (0, 5, 10), distribution focal loss weight dfl (1.5, 2.0), backbone freeze depth (0, 10), and label smoothing (0.0, 0.1). The production recipe confirmed here would be directly applied to the P1-P15 training: a higher-quality second model version meant the annotation partner at USZ would need fewer manual corrections when processing patients P16-P53. Rotation degrees was the dominant factor: degrees=5 consistently improved R_Normal over degrees=0 across all freeze and dfl combinations. Mast cells present in all orientations in aspirate smears, and moderate rotation augmentation directly targets this. The dfl parameter showed a consistent directional effect: dfl=2.0 regressed R_Normal relative to dfl=1.5 across all other parameter combinations, suggesting that increasing localisation loss weight tilts model capacity away from the classification discrimination needed for the rare Normal class. Freeze depth confirmed the Phase 2 finding: freeze=10 outperformed freeze=0 in every recall-relevant comparison. Label smoothing proved inert in this Ultralytics version, likely because the implementation does not propagate the parameter through the classification loss path.

The production recipe confirmed by the tinkering matrix was: degrees=5, dfl=1.5, freeze=10, lr0=0.001, cls=1.0, augment=True, mosaic=0.0, flipud=0.5.

An automated hyperparameter search using YOLO's built-in model.tune() was run on the production recipe (30 iterations, 50 epochs, patient P4 as holdout, mAP50-95 as fitness). The tuner converged on degrees=0.0 and dfl=2.47, both directions that the tinkering matrix had identified as harmful to R_Normal. The explanation is structural: a single-fold tuner optimizing mAP50-95 on a P4 holdout (an Atypisch-dominant patient) maximizes Atypisch-class performance at the expense of Normal-class generalization. Single-fold automated tuning is blind to cross-patient per-class recall balance. The production recipe from the tinkering matrix was retained.

A model-size comparison (P3-E) substituted yolo11l (43 M parameters) for yolo11n under identical conditions on the P1-P8 corpus and produced R_Normal = 0.695, the lowest value recorded across the entire project (mechanism in §Model Selection and Architecture).

The production recipe confirmed by Phase 3 was carried forward and applied to train on the full P1-P15 corpus, producing the second model version (DL_Modell_FV_v2.pt) that the annotation partner at USZ used for patients P16-P53.

Phase 4: Corpus expansion to 15 patients.

Phase 4 expanded the training corpus to 15 patients (1'564 images, 1'288 Atypisch, 391 Normal, 2'464 background images including FP; raw ratio 3.3:1, effective ratio 1.5:1 after automatic per-patient oversampling). The most significant addition was patient P9, which contributed 195 Normal annotations, more than doubling the total Normal count relative to the P1-P8 corpus. A 15-fold LOPO protocol was applied.

The first Phase 4 run (P4-A) was configured with `degrees=0` due to an omitted configuration variable; the second run (P4-B) applied the full production recipe including `degrees=5`. The delta between the two runs was within the cross-fold standard deviation on all metrics, indicating that the benefit of `degrees=5` observed at P1-P8 scale either diminishes with a more diverse 15-patient corpus or is within measurement noise at this scale. A freeze ablation run (P4-C, `freeze=0`) produced results essentially equivalent to P4-B across all metrics, confirming that backbone freezing remains appropriate and the P3-B recipe transfers to the larger corpus without re-tuning. Full per-fold results for P4-A, P4-B, and P4-C are reported in the Results section.

Phase 5: Full corpus (P1-P53).

The Phase 5 training run (P5-A) applied the confirmed production recipe (`degrees=5`, `df=1.5`, `freeze=10`, `lr=0.001`, `cls=1.0`) with fully automatic per-patient oversampling targeting a 2:1 Atypisch:Normal ratio across a LOPO design over all 35 annotated patients.

Model Testing

Performance estimates in this thesis are reported under two different evaluation protocols depending on the phase. Phase 1 and Phase 2 used stratified k-fold cross-validation (5 folds), where images were split randomly across folds with class stratification. From Phase 3 onward, the evaluation protocol was changed to LOPO, where all images from one patient are withheld from training entirely and used exclusively for evaluation. The switch was motivated by a systematic bias identified in the Phase 1 and Phase 2 results: stratified k-fold splits on small medical imaging datasets allow patient-level texture and staining style to leak from training into validation, artificially improving validation metrics without reflecting true cross-patient generalization. LOPO results and stratified k-fold results are therefore not directly comparable and are reported separately in the Results section.

Choosing LOPO over stratified k-fold directly addresses a gap identified in the broader literature. Ghete et al. [1] found that performance metrics across published bone marrow cell classification studies are not directly comparable in part because most studies evaluate on subsets of their own data, with few reporting results on patient-held-out splits and fewer still validating externally. LOPO is the strongest within-dataset approximation of external validation available without access to a second institution's data: each fold evaluates the model on a patient it has never seen in any form, under the same conditions a deployed system would face in routine clinical use. The cross-fold standard deviation reported in this thesis therefore reflects genuine inter-patient biological variability rather than statistical noise from random sampling, making the performance estimates more conservative and more informative than single-split or random k-fold results typically reported in the field.

Each fold produced a training split, a validation split used only for early stopping during training, and a test split consisting of the held-out patient's images. The metrics reported throughout this thesis refer to the test split. Across folds, the mean and standard deviation of each metric were computed. The cross-fold standard deviation reflects genuine inter-patient biological variability; patients with very small annotated sets (for example, P13 with 7 images) produce test metrics with extreme sensitivity, where a single missed cell changes recall by more than 14 percentage points.

A per-class recall extraction issue affected single-class holdout folds through Phase 4; affected folds are excluded from per-class mean calculations throughout. The issue is fully described, with the fix applied in Phase 5, in Section .

Evaluation Metrics

Model performance was assessed using standard metrics from object detection. Each detection outcome is first classified into one of four categories based on whether the model’s predicted bounding box overlaps sufficiently with a ground-truth annotation (at a minimum Intersection over Union of 0.50):

- **True Positive (TP)**: the model correctly detected a mast cell and the predicted box sufficiently overlapped the ground-truth annotation.
- **True Negative (TN)**: the model correctly produced no detection in a tile containing no annotated mast cell.
- **False Positive (FP)**: the model predicted a mast cell where the ground truth showed none (a false alarm).
- **False Negative (FN)**: the model failed to detect a mast cell that was present in the ground truth (a missed cell).

From these, the following performance indicators were computed [PA2_FV]:

Precision measures the proportion of predicted detections that were correct:

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall measures the proportion of ground-truth mast cells that the model successfully detected:

$$\text{Recall} = \frac{TP}{TP + FN}$$

mAP50 (Mean Average Precision at IoU 0.50) is the mean of the per-class Average Precision computed at a bounding-box overlap threshold of 50%. AP is the area under the precision-recall curve; mAP50 averages this across all classes:

$$\text{mAP50} = \frac{1}{N} \sum_{i=1}^N AP_{i@50}$$

mAP50-95 extends mAP50 by computing AP (Average Precision) at ten evenly spaced IoU thresholds from 0.50 to 0.95 (`torch.linspace(0.5, 0.95, 10)`) and averaging the result across thresholds and classes [4]. It penalises imprecise bounding box localisation more heavily than mAP50.

Recall is the primary metric for this project. A missed mast cell (false negative) causes the Atypisch fraction to be underestimated, potentially placing a patient below the WHO 25% diagnostic threshold and delaying diagnosis of SM. A false positive is less consequential for two reasons: the clinical workflow includes a manual review step where false alarms can be dismissed, and the diagnostic impact of a false positive depends on how far the true ratio is from the 25% threshold. In a patient with a strongly atypical cell population, a moderate false positive rate does not change the clinical conclusion because the Atypisch fraction remains clearly above the threshold. False

positives become clinically relevant primarily in borderline cases where the ratio is close to 25%, and it is precisely in those cases that the manual review step is most likely to be exercised.

Two per-class recall metrics are tracked separately throughout the experiments: `R_Atypisch` (recall for the Atypisch class) and `R_Normal` (recall for the Normal class). Both are clinically necessary because the WHO criterion depends on the ratio of Atypisch to total detected mast cells. Undercounting Normal cells inflates this ratio, potentially triggering a false-positive SM classification in borderline patients. Tracking both classes independently, rather than reporting a single aggregate recall, is a direct consequence of aligning the evaluation to the diagnostic question rather than to standard object detection benchmarking practice. Published studies in adjacent domains, such as general white blood cell subtype classification, typically optimize for aggregate mAP or F1 across all classes [3]; this thesis treats `R_Normal` as a primary metric in its own right because the WHO threshold makes the Normal count as diagnostically load-bearing as the Atypisch count. `R_Normal` received particular experimental attention because Normal is the minority class in the training corpus, consistently exhibiting lower recall and higher cross-fold variability than `R_Atypisch`.

mAP50 is used as the standard benchmark for run-to-run comparison across the experimental timeline, providing a single-number summary that captures both detection and classification quality. mAP50-95 is reported where available but assigned less interpretive weight; localisation precision at high IoU thresholds is clinically less relevant than detection recall for this task. Precision is monitored throughout but treated as a secondary metric.

UI Creation

Gradio Prototype

An initial web-based interface was developed using Gradio to support early-stage result visualization and model evaluation during development. The prototype provided two functional tabs. The Analysis tab ran batch inference on image directories and displayed tiles with YOLO detection overlays (Figure 1), producing labeled output images. The Research tab supported ground-truth versus prediction comparison (Figure 2, Figure 3), applying a lower confidence threshold to maximize recall sensitivity during evaluation. A statistics dashboard (Figure 4) aggregated per-class detection counts and displayed the Atypisch fraction relative to the WHO 25% threshold. The Gradio prototype served its purpose as a rapid development tool and was later superseded by the desktop application.




Figure 1: Analysis tab of the Gradio prototype showing uploaded image tiles with YOLO detection overlays.




Figure 2: Research tab showing ground-truth annotations alongside model predictions for direct visual comparison.




Figure 3: Research tab metrics output: precision-recall curves and performance plots generated from the ground-truth evaluation run.




Figure 4: Statistics dashboard of the Gradio prototype showing per-class detection counts and the Atypisch fraction relative to the WHO 25% diagnostic threshold.

Desktop Application

As the requirements for clinical deployment became clearer, it became apparent that a browser-based interface was not adequate for the actual usage context at USZ. A full patient slide scan produces approximately 20,000 image tiles, and a Gradio-based interface was not designed for streaming inference at that scale. Cinderella, the cell classification software already in use at USZ, was considered as an integration target, but embedding the pipeline into an existing hospital codebase within the thesis timeline was not feasible: the time required to understand a production clinical codebase combined with the security validation and change management processes required for hospital software integration extended well beyond the project scope.

A native desktop application (MastCellDetector) was developed as the primary clinical deliverable. The application was built with PyQt and packaged using PyInstaller into a self-contained executable that requires no Python environment on the target machine. While the primary deployment target is Windows workstations at USZ, PyInstaller equally supports Linux builds, making the same application distributable on Linux servers or HPC systems.

The intended workflow for a USZ analyst proceeds as follows. The analyst opens a local folder containing a patient's image tiles. The sidebar provides controls for confidence threshold, IoU non-maximum suppression threshold, and hardware selection (GPU with automatic VRAM-based batch size recommendation, or CPU fallback). Running inference streams YOLO-format label files to disk alongside each image in real time, so results are preserved if the run is interrupted.

The gallery view (Figure 5) displays all images with filter tabs for All images, With detections (Figure 6), No detection, Low confidence, Atypisch only, and Normal only. This allows a haematologist to navigate directly to the cases most likely to need manual review rather than scrolling through tens of thousands of tiles.




Figure 5: Gallery view of the desktop application showing all image tiles for a patient folder with sidebar controls for inference configuration.




Figure 6: Gallery filtered to tiles with detections, allowing rapid navigation to images where the model found mast cells.

An annotation editor is embedded in the gallery. Clicking any thumbnail opens the image with interactive bounding boxes (Figure 7): boxes can be repositioned by dragging, resized by dragging corner handles, and deleted with the Delete key. New boxes are drawn by pressing A and clicking and dragging a region, with class selected from a dropdown. Changes are saved to the corresponding label file with Ctrl+S or automatically on navigation to the next image. When a box is deleted as a false positive, an empty label file is written, marking the tile as a confirmed negative for any future retraining cycle. Edited files are protected from being overwritten by subsequent inference runs.




Figure 7: Annotation editor with interactive bounding boxes; boxes can be added, moved, resized, or deleted directly in the application without external tooling.

The Statistics tab (Figure 8) aggregates per-image detection counts into a patient-level summary that displays the total detected mast cells, the Atypisch fraction as a percentage, and a visual marker relative to the WHO 25% diagnostic threshold. This summary can be exported for clinical documentation or further analysis.




Figure 8: Statistics tab aggregating per-image detections into a patient-level summary with the Atypisch fraction displayed against the WHO 25% diagnostic threshold.

Implementation in USZ Infrastructure

The primary deliverable of this project is the MastCellDetector desktop application, distributed as a self-contained Windows directory. The application operates entirely offline after installation, which satisfies a practical requirement for deployment in hospital network environments where internet access from clinical workstations is restricted or prohibited.

The model shipped with the final version of the desktop application is the P6-A deployment model, trained on the complete P1-P53 corpus minus P11 using the production recipe established across Phases 3 and 4.

The final deployment model (P6-A) was trained on 1,832 images across 34 annotated patients (P1–P53 minus P11), achieving val mAP50-95 = 0.842, Precision = 0.906, Recall = 0.921 at best.pt (early stopping at epoch 657, best epoch 607). Weights: [hpc/tinker_final_holdoutP11/final/weights/best.pt](#).

The annotation editing and label export functionality in the application also addresses a secondary deliverable: the confirmed false positives and corrected labels produced during clinical use feed directly back into the training pipeline as the next generation of hard-negative examples, continuing the iterative annotation strategy established from Phase 1 onward.

The Linux build capability opens a second deployment path beyond the clinical workstation. A Linux build of MastCellDetector can be deployed on a computational server or HPC node at USZ, allowing the tool to be integrated directly into the institution’s scan-to-inference-to-evaluation pipeline: whole-slide images are scanned, tiles are passed through the model at server-side GPU throughput, and the resulting label files and statistics export are produced without requiring a Windows workstation in the loop. This makes the same application suitable for both interactive

clinical review and automated batch processing, depending on the deployment context. The source code of the application is publicly available [5].

3. Results

The following training hyperparameters are referenced throughout this section:

- **lr0** — initial learning rate. Controls the step size at the start of fine-tuning; too high a value causes catastrophic forgetting of pretrained weights.
- **freeze** — number of backbone layers frozen during training. **freeze=10** locks the YOLO11n backbone and trains only the detection head; **freeze=0** allows full fine-tuning.
- **degrees** — maximum rotation angle for on-the-fly augmentation. Increases orientation diversity in the training distribution.
- **dfl** — Distribution Focal Loss weight. Scales the bounding-box regression loss relative to the classification loss; higher values shift the model toward localisation accuracy.
- **cls** — classification loss weight. Scales the class-prediction loss; used here to compensate for class imbalance between Atypisch and Normal.
- **mosaic** — probability of mosaic augmentation (four tiles composited into one image). Disabled (**mosaic=0.0**) throughout because mosaic scrambles single-cell 512 px crops.
- **flipud / fliplr** — vertical and horizontal flip augmentation probabilities.
- **FP_NEG_OVERSAMPLE** — repeat factor for confirmed false-positive tiles in the training split. Values above 1 up-weight hard negatives.
- **BG_RATIO** — ratio of unannotated background tiles added as negatives per training image.

3.1. Model Results and Statistical Evaluation

Performance metrics for Phase 3 and Phase 4 runs are reported from the test split of each LOPO fold. Phase 1 and Phase 2 results derive from held-out validation folds under 5-fold stratified cross-validation and are not directly comparable to Phase 3–4 LOPO estimates, for the reasons described in the Model Testing section. All values quoted in this section refer to cross-fold means unless stated otherwise.

Phase 1. The P1-baseline achieved validation $\text{mAP}_{50-95} = 0.649 \pm 0.131$, precision = 0.768 ± 0.117 , and recall = 0.719 ± 0.178 . Cross-fold variance was high: Fold 1 reached $\text{mAP}_{50-95} = 0.805$ while Fold 5 collapsed at epoch 29 (recall = 0.432). No per-class breakdown was available as P1 contained only Atypisch annotations.

Phase 2. The 3×4 factorial DoE grid (11 cells, FP oversampling $\in \{1\times, 2\times, 3\times\}$ crossed with background ratio $\in \{0, 1, 2, 3\}$) found no statistically distinguishable difference across compositions: cell-level mAP_{50-95} variation was smaller than the within-cell cross-fold standard deviation. Background tiles decreased performance at every FP oversampling level. The fp=1, bgr=0 configuration (P2-H) was the grid winner at every metric, achieving $\text{mAP}_{50-95} = 0.494 \pm 0.166$, precision = 0.934 ± 0.057 , and recall = 0.537 ± 0.210 .

Phase 3. The structural pivot to 8-fold LOPO on 8 patients with on-the-fly per-epoch augmentation produced the largest single performance step in the project. P3-A achieved $\text{mAP}_{50-95} = 0.718 \pm 0.137$, precision = 0.764 ± 0.138 , and recall = 0.840 ± 0.122 , a 30.3 percentage point improvement in recall over P2-H (0.537). The cross-fold standard deviation dropped from 0.166 to 0.137 despite the harder evaluation protocol.

The tinkering matrix (P3-B, 14 of 24 cells completed) identified **degrees=5** rotation augmentation as the dominant training-time lever. The winning configuration (**degrees=5**, **dfl=1.5**, **freeze=10**, **label_smoothing=0.0**) achieved $\text{mAP}_{50-95} = 0.741 \pm 0.122$, precision = $0.812 \pm$

0.153, recall = 0.866 ± 0.131 , R_Atypisch = 0.892 ± 0.173 , and R_Normal = 0.840 ± 0.176 . `dfl=2.0` consistently regressed R_Normal relative to `dfl=1.5` across all degree and freeze combinations; `freeze=0` underperformed `freeze=10` in every recall-relevant slice; label smoothing was empirically inert (bit-identical fold metrics at `ls=0.0` and `ls=0.1`). These findings defined the production recipe (`degrees=5`, `dfl=1.5`, `freeze=10`, `lr0=0.001`, `cls=1.0`) adopted for all subsequent phases.

The automated hyperparameter tuner (P3-C, 30 iterations $\times$ 50 epochs, single P4 holdout, mAP50-95 fitness) converged on `degrees=0.0` and `dfl=2.47`. The production recipe from P3-B was retained.

The model-size comparison (P3-E, yolo11l, 43 M parameters) produced mAP50-95 = 0.745 ± 0.095 , precision = 0.778 ± 0.134 , recall = 0.785 ± 0.103 , and R_Normal = 0.695 ± 0.144 , the lowest R_Normal across all runs and 14.5 percentage points below P3-B.

Phase 4. Adding patients P9–P15 (424 new images, 293 new Normal annotations) expanded the corpus to 1,564 images and 15 LOPO folds. P4-A (`degrees=0`, the production recipe with the rotation variable inadvertently unset) achieved mAP50-95 = 0.793 ± 0.106 , precision = 0.851 ± 0.097 , recall = 0.904 , R_Atypisch = 0.913 , and R_Normal = 0.890 . The R_Normal gain from corpus expansion (+9.5 pp over P3-A) exceeded the largest single hyperparameter gain observed in Phase 3 (+4.5 pp from `degrees=5` in P3-B).

P4-B (full production recipe with `degrees=5` explicitly set) produced mAP50-95 = 0.789 ± 0.118 , precision = 0.826 ± 0.131 , recall = 0.895 ± 0.132 , R_Atypisch = 0.900 ± 0.174 , and R_Normal = 0.882 ± 0.193 . All P4-A vs P4-B deltas were within 1.6 percentage points. The freeze ablation (P4-C, `freeze=0`) produced mAP50-95 = 0.782 ± 0.124 , precision = 0.863 ± 0.087 , recall = 0.893 ± 0.135 , R_Atypisch = 0.896 ± 0.172 , and R_Normal = 0.882 ± 0.193 , all within 0.4 percentage points of P4-B. P4-B is the primary reported result for the P1–P15 phase, as all training parameters are explicitly set.

Phase 5. Phase 5 evaluated the pipeline at full deployment scale across all 35 annotated patients from the 53-patient corpus. Two protocols were applied: a 35-fold LOPO run (P5-A) and a grouped 5-fold CV (P5-B). P5-B is the primary generalization result; P5-A is an optimistic upper bound documented for methodological completeness.

P5-A: 35-fold LOPO (optimistic upper bound). The production recipe was applied across all 35 annotated patients. Of the 35 folds, 34 completed within the 48-hour SLURM allocation; Fold 17 (P37 holdout) was retrained independently and evaluated in a post-hoc inference pass (see Section).

Across 35 folds, post-hoc evaluation reported mAP50-95 = 0.858 ± 0.120 , precision = 0.925 ± 0.128 , recall = 0.951 ± 0.084 , R_Atypisch = 0.920 ± 0.129 ($n=20$ folds with Atypisch ground truth), and R_Normal = 0.941 ± 0.119 ($n=29$ folds with Normal ground truth). Of the 35 folds, 19 reported mAP50 at the metric ceiling (0.995), corresponding to patients from the P16–P53 acquisition batch who contributed only one to three annotated images each. On a one-image test set, a single correct detection produces recall = 1.000 regardless of model generalization capacity, inflating the aggregate mean.

P5-B: Grouped 5-fold CV (primary Phase 5 result). The 35 annotated patients were partitioned into five fixed groups, each serving as the test holdout exactly once. Every group contains at least one Atypisch-annotated and one Normal-annotated patient; the two largest patients (P1:

428 images, P2: 417 images) each anchor a separate group; test-image counts are balanced across groups (315–443 images per group). All five folds used the confirmed production recipe.

Across five groups, $mAP_{50-95} = 0.748 \pm 0.079$, precision = 0.864 ± 0.022 , recall = 0.853 ± 0.087 , $R_{\text{Atypisch}} = 0.885 \pm 0.035$, and $R_{\text{Normal}} = 0.821 \pm 0.165$. Per-group metrics are in Section . G3 achieved $R_{\text{Normal}} = 0.981$, the highest across all multi-patient evaluation folds in the project.

Phase 6. A G1–G4 ablation run re-ran four-fold grouped CV on G1–G4 only, with G5 excluded from training, validation, and negative sets. Results against the corresponding P5-B G1–G4 folds: $mAP_{50-95} = 0.738$ (P5-B G1–G4: 0.740), Recall = 0.849 (+0.006), $R_{\text{Atypisch}} = 0.897$ (+0.009), $R_{\text{Normal}} = 0.801$ (+0.004). All deltas fall within cross-fold noise, confirming that training without G5 is feasible with no substantial performance loss.

This finding raised the question of whether G5 could be withheld entirely as a calibration cohort. That approach was ruled out: to make use of all G5 patients for calibration, inference would need to be run across the full slide set of each G5 patient (roughly 20,000 tiles per patient, approximately 380'000 tiles in total), which is impractical for a single calibration pass. The alternative, withholding all of G5 from training but calibrating on only one of its patients, would discard image information from the remaining G5 patients without benefit. Instead, a single patient was withheld from training: P11, selected because its 17 Atypisch and 20 Normal annotations provide near-equal class balance, the most informative single-patient confidence reference in the corpus. All remaining patients, including the rest of G5, were retained in training. The final deployment model was trained on P1–P53 minus P11 (1,832 images, 1,459 Atypisch, 491 Normal, 3.0:1 ratio) using the production recipe unchanged.

Table 2 summarizes the key milestone runs across all phases. The complete experiment record, including corrupted runs and all 11 DoE grid cells, is maintained in [hpc/experiment_log.md](#) [6]. P3-A per-class recall values in Table 2 are cross-fold means from the tinkering matrix reference configuration (degrees=0, dfl=1.5, freeze=10, label_smoothing=0.0), which is equivalent to the direct P3-A run; folds where the respective class was absent from the held-out patient are excluded from each per-class mean.

Table 2: Milestone run progression summary across all experimental phases. All metrics shown as mean $\pm$ std across folds except P6-A (single training run, val split only; no held-out test set). Base model is yolo11n throughout except P3-E (yolo11l, 43 M parameters). CV-codes: S = stratified k-fold, L = leave-one-patient-out (LOPO), CV = grouped cross-validation; prefix = fold/group count. *P3-A per-class recall values are cross-fold means from the tinkering matrix reference configuration (degrees=0, dfl=1.5, freeze=10); folds where the respective class was absent from the held-out test set are excluded from each per-class average.

Run	CV	mAP50-95	Prec.	Recall	R_Atyp	R_Norm
P1-base	5f-S	0.649 $\pm$ 0.131	0.768 $\pm$ 0.117	0.719 $\pm$ 0.178	—	—
P2-C	5f-S	0.344 $\pm$ 0.031	0.880 $\pm$ 0.017	0.397 $\pm$ 0.022	—	—
P2-H	5f-S	0.494 $\pm$ 0.166	0.934 $\pm$ 0.057	0.537 $\pm$ 0.210	—	—
P3-A	8f-L	0.718 $\pm$ 0.137	0.764 $\pm$ 0.138	0.840 $\pm$ 0.122	0.884 $\pm$ 0.170*	0.795 $\pm$ 0.164*
P3-B	8f-L	0.741 $\pm$ 0.122	0.812 $\pm$ 0.153	0.866 $\pm$ 0.131	0.892 $\pm$ 0.173	0.840 $\pm$ 0.176
P3-E	8f-L	0.745 $\pm$ 0.095	0.778 $\pm$ 0.134	0.785 $\pm$ 0.103	0.875 $\pm$ 0.169	0.695 $\pm$ 0.144
P4-A	15f-L	0.793 $\pm$ 0.106	0.851 $\pm$ 0.097	0.904 $\pm$ 0.120	0.913 $\pm$ 0.110	0.890 $\pm$ 0.194
P4-B	15f-L	0.789 $\pm$ 0.118	0.826 $\pm$ 0.131	0.895 $\pm$ 0.132	0.900 $\pm$ 0.174	0.882 $\pm$ 0.193
P4-C	15f-L	0.782 $\pm$ 0.124	0.863 $\pm$ 0.087	0.893 $\pm$ 0.135	0.896 $\pm$ 0.172	0.882 $\pm$ 0.193
P5-A	35f-L	0.858 $\pm$ 0.120	0.925 $\pm$ 0.128	0.951 $\pm$ 0.084	0.920 $\pm$ 0.129	0.941 $\pm$ 0.119
P5-B	5g-CV	0.748 $\pm$ 0.079	0.864 $\pm$ 0.022	0.853 $\pm$ 0.087	0.885 $\pm$ 0.035	0.821 $\pm$ 0.165
P6-A	val	0.842	0.906	0.921	0.914	0.929

P1-base established the initial proof of concept on a single patient. P2-C identified the learning rate mismatch that suppressed recall in all DL_Model_FV runs before the fix; P2-H is the Phase 2 reference after the learning rate and freeze correction. P3-A marks the single largest structural improvement in the project. P3-B confirmed the production recipe through the 8-fold tinkering matrix and is the first fully validated configuration (bold). P3-E is a deliberate negative result bounding model-size choices. P4-A documents the corpus expansion gain from P9–P15; P4-B is the primary result for the P1–P15 phase with all recipe parameters explicitly set. P4-C is the freeze ablation. P5-A is the optimistic upper bound at full corpus scale; P5-B is the primary cross-validated generalization result. P6-A is the final deployment model trained on all 35 annotated patients minus P11; its metrics are validation-split values as no held-out test set exists.

3.2. Quantitative and Qualitative Evaluation

Per-Fold Breakdown (P5-A)

Table 3 presents the complete 35-fold post-hoc evaluation results for P5-A. Each row corresponds to one annotated patient held out from training entirely; all metrics are evaluated on that patient’s images exclusively. Post-hoc evaluation was applied uniformly across all 35 folds because Fold 17 (P37 holdout) was killed by the 48-hour SLURM wall-time limit and retrained independently; running the unified post-hoc pass over all folds ensures a consistent result set without mixed evaluation protocols (see Section). Patients are listed in numeric order.

Table 3: P5-A per-fold test results (35-fold LOPO post-hoc evaluation, full P1–P53 annotated corpus). GT Atyp / GT Norm = ground-truth cell instance counts in the held-out test split. “—” indicates the class was absent from that test split and is excluded from the per-class mean. R_Atyp mean over n=20 folds; R_Norm mean over n=29 folds.

Patient	GT Atyp	GT Norm	mAP50-95	mAP50	Prec.	Recall	R_Atyp	R_Norm
P1	482	2	0.610	0.813	0.577	0.973	0.946	1.000
P2	449	6	0.600	0.716	0.721	0.712	0.924	0.500
P3	60	4	0.736	0.906	0.785	0.931	0.911	0.951
P4	122	6	0.675	0.806	0.754	0.796	0.926	0.667
P5	3	18	0.652	0.760	0.595	0.740	0.667	0.814
P6	0	6	0.846	0.995	1.000	1.000	—	1.000
P7	8	12	0.809	0.975	0.929	0.917	1.000	0.833
P8	0	24	0.816	0.986	0.914	0.958	—	0.958
P9	4	195	0.600	0.622	0.635	0.722	0.500	0.944
P10	5	20	0.909	0.991	0.987	0.906	1.000	0.812
P11	17	20	0.863	0.964	0.968	0.887	0.941	0.832
P12	133	0	0.832	0.984	0.962	0.952	0.952	—
P13	0	7	0.920	0.995	1.000	1.000	—	1.000
P14	5	3	0.933	0.995	0.978	1.000	1.000	1.000
P15	0	68	0.949	0.991	0.984	0.985	—	0.985
P16	22	1	0.834	0.994	1.000	0.977	0.955	1.000
P17	69	0	0.832	0.992	0.945	0.994	0.994	—
P18	81	1	0.948	0.993	0.934	0.946	0.893	1.000
P24	1	0	0.796	0.995	1.000	1.000	1.000	—
P28	1	0	0.995	0.995	1.000	1.000	1.000	—
P29	1	0	0.896	0.995	1.000	1.000	1.000	—
P31	0	1	0.995	0.995	1.000	1.000	—	1.000
P37	0	2	0.896	0.995	1.000	1.000	—	1.000
P39	0	1	0.896	0.995	1.000	1.000	—	1.000
P40	0	1	0.995	0.995	1.000	1.000	—	1.000
P43	0	3	0.940	0.995	1.000	1.000	—	1.000
P44	0	2	0.921	0.995	1.000	1.000	—	1.000
P45	0	1	0.995	0.995	1.000	1.000	—	1.000
P46	0	4	0.967	0.995	1.000	1.000	—	1.000

Patient	GT Atyp	GT Norm	mAP50-95	mAP50	Prec.	Recall	R_Atyp	R_Norm
P47	0	3	0.918	0.995	1.000	1.000	—	1.000
P48	7	0	0.857	0.995	1.000	1.000	1.000	—
P50	0	1	0.995	0.995	1.000	1.000	—	1.000
P51	5	90	0.700	0.765	0.740	0.900	0.800	1.000
P52	1	4	0.945	0.995	0.971	1.000	1.000	1.000
P53	0	5	0.952	0.995	1.000	1.000	—	1.000
Mean			0.858	0.948	0.925	0.951	0.920 (n=20)	0.941 (n=29)

P1 R_Normal = 1.000 based on 2 Normal ground-truth instances and is not a stable estimate. P12 has 133 Atypisch and 0 Normal ground-truth instances; R_Normal is absent from its test set.

Nineteen folds reported mAP50 at the ceiling value of 0.995. The majority are patients from the P16–P53 acquisition batch who contributed one to three annotated images each; P6, P13, and P14 also reach the ceiling for the same structural reason. On a one or two-image test set a single correct detection drives recall and mAP50 to their respective maxima regardless of model generalization capacity. These folds confirm that the model produces no catastrophic failures on any individual patient from the expanded corpus, but they carry no cross-patient generalization signal. The five patients with mAP50 below 0.820 (P2, P4, P5, P9, P51) represent the meaningful lower tail of the performance distribution at full corpus scale; their per-fold patterns are examined in Section .

Per-Group Breakdown (P5-B)

Table 4 presents the complete per-group test results for P5-B. Each row corresponds to one held-out patient group; all metrics are evaluated exclusively on images from patients in that group. The holdout column identifies the constituent patients; G5 contains the six large-corpus patients from P4 and P5 plus all 13 remaining small-corpus patients from P16–P53.

Table 4: P5-B grouped 5-fold CV split composition and test results (35 annotated patients, P1–P53 corpus). Train = annotated (positive) images; each fold additionally receives 3,773-4,595 negative tiles. Atyp/Norm = ground-truth label counts in the test split. All five groups have both classes present.

Group	Holdout patients	Train	Val	Test	Atyp	Norm	mAP50-95	mAP50	Prec.	Recall	R_Atyp	R_Norm
G1	P1, P13, P31	1219	214	436	482	10	0.649	0.766	0.850	0.746	0.904	0.588
G2	P2, P8, P37	1213	213	443	449	32	0.679	0.858	0.839	0.773	0.836	0.710
G3	P9, P12, P53	1315	231	323	137	200	0.813	0.913	0.882	0.932	0.884	0.981
G4	P3, P6, P15–P18, P48	1322	232	315	239	80	0.821	0.918	0.892	0.919	0.928	0.909
G5	P4, P5, P7, P10–P11, P14, P24, P28–P29, P39–P40, P43–P47, P50–P52	1289	228	352	169	189	0.777	0.905	0.854	0.896	0.873	0.918
Mean ± std							0.748 ± 0.079	0.872 ± 0.064	0.864 ± 0.023	0.853 ± 0.087	0.885 ± 0.035	0.821 ± 0.165

3.3. Evaluation Visualization

Training convergence for P5-B is shown in Figure 9 for a representative fold (Fold 3, G3 holdout: P9, P12, P53; mAP50 = 0.913, recall = 0.932). The training and validation loss curves converge within approximately 150 epochs on the majority of folds, with YOLO’s early-stopping criterion (patience = 50) triggering at the validation mAP50 plateau.




Figure 9: Training convergence for P5-B Fold 3 (G3 holdout: P9, P12, P53). Panels from top left: box loss, classification loss, and distribution focal loss for training and validation splits; precision, recall, mAP50, and mAP50-95 across epochs.

The normalized confusion matrix for the same fold is shown in Figure 10. Row normalization is applied because the Atypisch and Normal test-set sizes differ, making raw counts incomparable across classes; row-normalized values represent the fraction of each ground-truth class assigned to each predicted class. The main diagonal records 0.96 for Atypisch and 0.88 for Normal, indicating the fraction of each class correctly detected. Cross-class misclassifications between Atypisch and Normal are negligible at 0.01. The dominant off-diagonal entry is 0.91 in the background-to-Atypisch cell, reflecting the proportion of background tiles assigned to the Atypisch class.




Figure 10: Normalized confusion matrix for P5-B Fold 3 (G3 holdout: P9, P12, P53). Rows are ground-truth classes; columns are predicted classes. Values are row-normalized fractions. The main diagonal entries are 0.96 (Atypisch) and 0.88 (Normal); cross-class error is 0.01; background-to-Atypisch is 0.91.

Phase 5–6 Validation Comparison

Figure 11 shows per-group R_{Atypisch} and R_{Normal} for all five P5-B folds alongside the cross-fold means.




Figure 11: P5-B per-group recall by class. Bars show R_{Atypisch} (blue) and R_{Normal} (orange) for each held-out patient group. Dashed horizontal lines mark the cross-fold means ($R_{\text{Atypisch}} = 0.885$, $R_{\text{Normal}} = 0.821$).

Figure 12 shows mAP50-95, Precision, and Recall for P5-A (35-fold LOPO, test split), P5-B (5-group CV, test split), and P6-A (final deployment model, validation split). P6-A has no held-out test set as it was trained on all annotated patients minus P11.




Figure 12: Performance comparison across P5-A (35-fold LOPO, test), P5-B (5-group CV, test), and P6-A (final model, val). Error bars show ± 1 std across folds, clipped at 1.0. P6-A is a single training run with no held-out test set.

Figure 13 compares all six evaluation metrics between P5-A and P5-B. Error bars show one standard deviation across folds for both evaluations; R_Atyp and R_Norm for P5-A are computed over the subset of folds where the respective class was present in the test set.




Figure 13: Metric comparison between P5-A (35-fold LOPO) and P5-B (grouped 5-fold CV). Error bars represent ± 1 standard deviation across folds. R_Atyp and R_Norm for P5-A are computed over the 20 and 29 folds respectively where the class was present in the test set.

Final Model Training Visualization

Training convergence for the P6-A deployment model is shown in Figure 14. The model trained on 1,832 images (P1–P53 minus P11) and converged at epoch 607 (early stopping, patience=50).




Figure 14: Training curves for the P6-A final deployment model. Panels show box loss, classification loss, and DFL loss for training and validation splits; precision, recall, mAP50, and mAP50-95 across epochs.

The normalized confusion matrix at best.pt is shown in Figure 15. Evaluated on the validation split (15% of training data).




Figure 15: Normalized confusion matrix for P6-A (best.pt, evaluated on the val split). Rows are ground-truth classes; columns are predicted classes.

3.4. Real World Usage

Confidence Threshold Calibration

The final deployment model (P6-A) was subsequently applied to the complete P11 tile set. The resulting confidence distribution is shown in Figure 16.




Figure 16: Confidence distribution of the final deployment model (P6-A) across the complete P11 tile set (~12'000 tiles; 17 Atypisch and 20 Normal annotated).

4. Discussion

4.1. Interpretation of Results

To our knowledge, this is the first deep learning pipeline for the specific diagnostic task of distinguishing atypical from normal mast cells in bone marrow aspirates, validated under patient-held-out cross-validation and aligned to the WHO 25% minor criterion for systemic mastocytosis. The systematic review by Ghete et al. [1] of AI-based bone marrow cell classification from 2019 to 2024 documents work on broad cell-type classification but none targeting the Atypisch-versus-Normal mast cell distinction that the WHO ratio depends on, and few studies validating on patient-held-out splits at all. The work is also delivered, not just measured: the MastCellDetector desktop application [5] runs the final model offline at USZ and reports the Atypisch fraction directly against the 25% threshold, turning the pipeline into a tool that participates in the diagnostic decision rather than a benchmark number.

The primary generalization estimate is the P5-B grouped 5-fold cross-validation: $mAP_{50-95} = 0.748$, precision = 0.864, recall = 0.853, $R_{\text{Atypisch}} = 0.885$, and $R_{\text{Normal}} = 0.821$ across 35 annotated patients, with no image from a held-out patient ever appearing in training within its fold. These are conservative, deployment-relevant numbers: the cross-fold standard deviation reflects genuine inter-patient biological variability rather than random-split noise, and the test groups (315 to 443 images each) are large enough to absorb individual missed detections without catastrophic recall swings. The 35-fold LOPO run (P5-A: $R_{\text{Atypisch}} = 0.920$, $R_{\text{Normal}} = 0.941$) is reported as an optimistic upper bound only. Its per-class means are inflated by 19 folds at the mAP_{50} ceiling of 0.995, where a single-image test set guarantees perfect recall once its one annotated cell is detected, independent of any generalization capacity. P5-B is the number a clinician should trust; P5-A bounds the best case.

The final deployment model (P6-A), trained on all 34 annotated patients minus P11, achieves $R_{\text{Atypisch}} = 0.914$ and $R_{\text{Normal}} = 0.929$ on its validation split, consistent with the P5-B grouped-CV estimate and confirming that the production recipe generalizes cleanly to the expanded training set. These are in-sample validation metrics, not held-out test results, but their alignment with P5-B is reassuring evidence against overfitting at deployment scale.

R_{Normal} is the load-bearing metric for this task, because undercounting Normal cells inflates the Atypisch fraction and can push a borderline patient across the 25% threshold. The single largest improvement in R_{Normal} at any point in the project came from corpus expansion, specifically the addition of patient P9 with its 195 Normal annotations, and it exceeded the gain from any individual hyperparameter change. This is not an incidental observation. The Phase 2 design-of-experiment grid varied false-positive oversampling and background-tile ratio across the full composition space of a fixed corpus and found no statistically distinguishable difference: the spread across all configurations was smaller than the cross-fold standard deviation within any single run. The controlled conclusion is that when class imbalance is structural rather than accidental, that is, when the minority class is rare because the data is sampled from disease-positive patients in whom that morphology is genuinely uncommon, reweighting and resampling the existing data cannot substitute for collecting more biologically diverse patients. For rare-event medical imaging this is a transferable lesson: corpus expansion, not hyperparameter tuning, is the primary lever for rare-class recall under structural imbalance. The trajectory bears this out, with R_{Normal} reaching 0.981 on the best multi-patient fold (G3) once the corpus had acquired sufficient Normal-class coverage.

The automated hyperparameter tuner (P3-C) reinforces this point from a different angle. YOLO’s built-in tuner, run for 30 iterations on a single Atypisch-dominant holdout patient (P4), converged on $\text{degrees}=0.0$ and $\text{dfl}=2.47$, both parameters that the cross-validated tinkering matrix had identified as harmful to R_Normal. The tuner behaved rationally given its objective, mAP50-95 on a P4 holdout, but that objective rewards Atypisch-class performance at the expense of Normal-class generalization. This finding generalizes: for rare-class tasks where class imbalance is structural, standard automated hyperparameter search requires multi-patient cross-validation to prevent the dominant class from steering the search. Single-fold tuning is blind to the recall balance that the diagnostic criterion depends on. The production recipe retained from the tinkering matrix was not a conservative choice; it was the only choice that a cross-validated evaluation could identify as correct.

A second methodological contribution is the bootstrapped, model-assisted annotation loop. The first patient was annotated by hand to train an initial detector; each subsequent model generation then served as the CVAT inference backend for the next batch of patients, and the false positives the annotator confirmed at each round became curated hard negatives for the following training cycle. This coupling of annotation and model development is, to our knowledge, not documented for mast cell detection, and it is what made a 53-patient corpus feasible within a single thesis: the Ghete et al. review notes that most published studies apply models to pre-existing datasets rather than building the annotation pipeline alongside the model. Here the clinical partner’s annotation capacity was extended by each model generation, and each generation’s hardest mistakes were fed forward as training signal rather than discarded.

Finally, the threshold question is handled honestly. Rather than shipping a single hard-coded confidence value, the project withheld patient P11 from all training as a near-balanced calibration cohort, ran the final model over its full tile set out of sample, and delivered the resulting confidence histogram to USZ so the hospital can choose its own operating point against its own recall and review-burden tradeoff. A deployed tool that hands the institution an out-of-sample calibration reference, rather than an in-sample number, is the appropriate form for a clinical decision that is ultimately the haematologist’s to make.

The metric range across the P5-B groups is wide but structured, and each extreme is explained by corpus composition rather than model instability. G1’s mAP50 of 0.766 is depressed by the long-standing P1 Normal fragility (P1’s old Pos_neg annotation regime and only 10 Normal test instances in the group), not by a Phase 5 regression. G2’s recall of 0.773 reflects that P2 and P8 are the two most Atypisch-dominant patients in the corpus, so the G2 test set is unusually sensitive to any single missed Atypisch cell. G3, G4, and G5 each exceed 0.89 recall, converging with the P4-B LOPO mean of 0.895 on the earlier 15-patient corpus. That convergence is itself a result: training on the expanded 30-patient set maintained performance while adding 20 new patients to the evaluation scope, with no distributional drift relative to the P1 to P15 results. **absatz stimmt noch nicht ganz**

The grouped CV design is not equivalent to LOPO, but it preserves the property that matters: no image from a holdout patient appears in its fold’s training set. It was also the tractable choice. Five folds cost roughly 36 hours at $\text{MAX_PARALLEL} = 2$, fitting the 96-hour HPC allocation that made a full 35-fold LOPO impractical to run as the primary protocol. Grouped patient CV is therefore the deployable approximation of LOPO at this corpus scale.

Fold-Level Structural Patterns

P9 is the weakest informative fold. P9 contributes 4 Atypisch and 195 Normal instances to the test set. A single missed Atypisch cell moves $R_Atypisch$ by 25 percentage points; P5-A recorded $R_Atypisch = 0.500$ (2 of 4 detected). The resulting mAP50 of 0.622 is the lowest across all 35 folds. P9’s extreme class asymmetry (4 Atypisch against 195 Normal) makes it the hardest fold by construction: the model trains on a broadly Atypisch-dominant corpus but must classify correctly in a patient where Normal cells outnumber Atypisch by nearly 50:1. The high R_Normal of 0.944 on the same fold confirms that the model handles the Normal-dominant distribution correctly when P9 is withheld; the weakness is specific to the four rare Atypisch instances.

P1 recovered from its Phase 3–4 structural weakness. The P1 holdout fold, which produced mAP50 = 0.584 under the 14-patient P4-B training set, improved to mAP50 = 0.813 under the 34-patient P5-A training set. The annotation-regime gap between P1’s original `Pos_neg` directory structure and the CVAT-assisted P2–P53 annotations remains in place, but the expanded corpus provides enough additional Atypisch-dominant training material to partially compensate for the distributional mismatch. $R_Atypisch = 0.946$ and $R_Normal = 1.000$ (on 2 Normal instances) are both higher than in any prior phase under this holdout.

P2 and P5 show small-sample minority-class fragility. The P2 holdout (449 Atypisch, 6 Normal) produced $R_Normal = 0.500$ (3 of 6 detected); a single additional missed Normal cell would reduce this to 0.333. P5 (3 Atypisch, 18 Normal) recorded $R_Atypisch = 0.667$ (2 of 3 detected); one missed Atypisch shifts the value by 33 percentage points. Both figures are dominated by small-sample noise rather than systematic model error. P5’s $R_Normal = 0.814$ and P2’s $R_Atypisch = 0.924$ confirm that each patient’s majority class is detected reliably; the fragility is confined to the minority class within each specific test set.

Ceiling folds are structurally uninformative. The 19 folds at mAP50 = 0.995 are patients who contributed one to three annotated images each, predominantly from the P16–P53 acquisition batch but including three early patients (P6, P13, P14). Perfect or near-perfect recall on a one-image test set is mathematically guaranteed when the single annotated cell is detected; these values confirm the absence of catastrophic failures but carry no information about cross-patient generalization capacity. Per-fold recall values throughout Table 3 are informative as lower bounds (for structurally fragile small-sample folds) and as ceiling observations rather than as stable point estimates.

Measurement Notes

A per-class recall extraction issue affected all LOPO runs through Phase 4. The utility function indexed Ultralytics’ per-class result arrays by position rather than by the `ap_class_index` field. When a holdout patient contained only one class and the model produced predictions for that class only, Ultralytics returned a one-element rather than a two-element result array, causing positional indexing to assign the single element to $R_Atypisch$ and report R_Normal as NaN. This affected Fold 5 (P13 holdout) in all Phase 3 and Phase 4 runs; the true R_Normal for that fold is 1.000, confirmed from the per-fold confusion matrix. Affected folds are excluded from per-class cross-fold means and clearly marked in all fold tables.

For the P5-A post-hoc evaluation, the bug was fixed by examining per-patient ground-truth label files before calling `val()`: the label-derived GT class counts provided a reliable class-index lookup independent of Ultralytics’ internal `ap_class_index` field. This fix allowed the 35-fold P5-A results to be extracted correctly without per-fold manual inspection; folds where only one class

had ground-truth instances were routed to the appropriate class slot by GT count rather than by position. The aggregate metrics mAP50 and overall Recall are computed internally by Ultralytics and are unaffected by this issue in all phases.

RQ1 asked which performance metrics are most critical for clinical reliability. The answer is per-class recall tracked independently for both morphological classes. A false negative, a missed mast cell, biases the Atypisch fraction at the diagnostic decision boundary; R_Normal is the harder of the two because Normal is the minority class in the training corpus. mAP50 serves as the run-to-run comparison summary. F1-score is not used as a primary metric because it treats false positives and false negatives symmetrically, whereas this task is explicitly asymmetric: false positives are recoverable at the manual review step, while false negatives silently distort the WHO ratio. Precision is monitored but secondary for the same reason.

RQ2 asked whether the pipeline achieves rare-event detection without unmanageable false positives. It does, with the caveat that operating threshold selection is left to the clinic. Mast cells are present in only a small fraction of tiles in a full slide scan (~20'000 tiles per patient), and the confirmed-false-positive hard-negative training loop suppresses the most common error modes across annotation generations. The conservative P5-B estimate achieves R_Normal = 0.821 at a false-positive rate that is manageable within the MastCellDetector review workflow: the representative G3 confusion matrix records main-diagonal entries of 0.96 (Atypisch) and 0.88 (Normal) with cross-class error at 0.01. The out-of-sample confidence histogram from patient P11 shows a bimodal distribution separating the false-positive bulk from the true-positive peak, giving the hospital a principled basis for threshold selection without inheriting a single hard-coded engineering value.

RQ3 asked whether the pipeline reliably distinguishes typical from atypical mast cell phenotypes to support the WHO >25% criterion. Yes. Under patient-held-out cross-validation the pipeline achieves R_Atypisch = 0.885 and R_Normal = 0.821 (P5-B, conservative estimate); the final deployment model reports R_Atypisch = 0.914 and R_Normal = 0.929 on its validation split (P6-A). Both classes that enter the Atypisch-to-total ratio are counted with reliable recall, which is the necessary condition for computing a trustworthy fraction. The MastCellDetector application directly operationalizes the answer: its Statistics tab aggregates per-tile detections into a patient-level Atypisch fraction and displays it against the 25% threshold, mapping the pipeline's output onto the diagnostic question without further manual computation.

4.2. Clinical Relevance

The clinical value of this pipeline follows directly from the diagnostic question it is aligned to. The WHO minor criterion for SM is met when at least 25% of mast cells in the aspirate are atypical [7], so the quantity that matters is a ratio of two counts, not a single detection score. The pipeline is built around that ratio: it counts Atypisch and Normal cells separately and reports the Atypisch fraction against the threshold.

This is why the error asymmetry is treated the way it is. A false negative, a missed mast cell, biases the ratio and can move a borderline patient across the 25% line and delay an SM diagnosis, whereas a false positive is recoverable at the manual review step and only matters near the threshold. The evaluation therefore prioritizes recall, and Normal-class recall in particular, because undercounting Normal cells inflates the Atypisch fraction toward a false-positive classification. Optimizing for aggregate mAP, as adjacent white-blood-cell studies typically do [3], would obscure exactly the per-class behaviour the diagnosis depends on.

The bridge from model to clinic is the MastCellDetector desktop application [5]. It runs entirely offline, which satisfies the network restrictions on hospital clinical workstations; it streams label files to disk during inference so a long run is recoverable; and it lets the haematologist filter tens of thousands of tiles down to the cases that need review (with detections, low confidence, single-class) instead of scrolling through a whole slide. Its Statistics tab produces the patient-level Atypisch fraction against the 25% threshold as an exportable summary for clinical documentation. The annotation editor closes the loop: corrections made during clinical use are written back as label files and confirmed negatives, feeding the next training generation. Rather than asking the clinic to trust a fixed threshold, the project delivered an out-of-sample confidence calibration from the withheld patient P11 so USZ sets its own operating point. The intended role is assistive: the tool does the counting and surfaces the candidates, and the haematologist makes the diagnosis.

4.3. Error Analysis

The pipeline’s errors cluster into four understood modes, none of which is a general detection failure.

Minority-class small-sample fragility. The largest visible errors are concentrated in folds where the held-out patient contributes only a handful of minority-class cells. In P5-A, P9 (4 Atypisch against 195 Normal) records $R_{\text{Atypisch}} = 0.500$ because two missed cells out of four move the metric by 50 points, and P2 (6 Normal) records $R_{\text{Normal}} = 0.500$ on three missed cells. These are properties of the test set, not of the model: each patient’s majority class is detected reliably on the same fold (P9 $R_{\text{Normal}} = 0.944$, P2 $R_{\text{Atypisch}} = 0.924$). Such folds are informative as lower bounds, not as point estimates, and are handled in the reported means by restricting per-class averages to folds with enough minority-class instances.

The P1 annotation-regime outlier. P1 was annotated under the original Pos_neg directory structure before the CVAT-assisted workflow, and its holdout fold was the weakest in Phase 3 and Phase 4 ($mAP50 = 0.584$ under the 14-patient training set). The expanded corpus partly compensated for this distribution mismatch (P1 $mAP50$ recovered to 0.813 under the 34-patient P5-A training set), but the annotation-regime difference remains a known source of error and is the main driver of G1’s depressed Normal recall in P5-B.

The per-class recall measurement bug. A positional-indexing error in the per-class extraction utility misreported R_{Normal} for single-class holdout folds (notably P13) through Phase 4; it is a measurement artefact, not a model error, and is documented in full in Section . Affected folds are excluded from per-class means, and the underlying detections are correct.

Out-of-distribution generalization. The most important limitation is structural: the training corpus is drawn almost entirely from confirmed or suspected SM patients, so the model has learned the mast-cell morphology distribution of an SM-enriched population. Its behaviour on aspirates from patients with unrelated haematological conditions, or from healthy controls, is out of distribution and not characterized by this work. The reported metrics describe performance within the SM diagnostic workflow they were built for, and should not be read as evidence of general bone-marrow cell-classification performance.

5. Conclusion

5.1. Summary of Key Findings

This thesis delivers, to our knowledge, the first patient-held-out-validated deep learning pipeline for distinguishing atypical from normal mast cells in bone marrow aspirates and aligning the result to the WHO 25% minor criterion for systemic mastocytosis. A YOLO11n detector trained on a 53-patient corpus (35 with annotated cells) generalizes across patients under grouped cross-validation at $\text{mAP}_{50-95} = 0.748$, $R_{\text{Atypisch}} = 0.885$, and $R_{\text{Normal}} = 0.821$, reaching $R_{\text{Normal}} = 0.981$ on the strongest multi-patient fold, with a 35-fold LOPO run bounding the optimistic case at $R_{\text{Atypisch}} = 0.920$ and $R_{\text{Normal}} = 0.941$. Three findings carry beyond the numbers. First, a bootstrapped model-assisted annotation loop, in which each model generation labels the next batch and its confirmed false positives become the next generation's hard negatives, made a 53-patient rare-cell corpus feasible within a single thesis. Second, a controlled design-of-experiment result showed that under structural class imbalance, corpus expansion improves rare-class recall more than any hyperparameter change, a transferable lesson for rare-event medical imaging. Third, the work ships a usable artefact: the MastCellDetector desktop application runs the model offline at USZ, reports the Atypisch fraction against the 25% threshold, and was accompanied by an out-of-sample confidence calibration so the hospital sets its own operating point.

5.2. Limitations and Challenges

The corpus is drawn almost entirely from SM-enriched patients, so performance on non-SM or healthy aspirates is out of distribution and uncharacterized. The minority Normal class remains the harder direction, and per-class recall on small single-patient folds is dominated by sampling noise rather than model behaviour, which is why the conservative grouped-CV estimate, not the ceiling-inflated LOPO mean, is treated as primary. No second-institution external validation set was available, so LOPO and grouped patient CV are the strongest generalization estimates obtainable from a single site. The per-class recall measurement bug, although corrected for in the reported means, affected single-class folds during development. Finally, the operating confidence threshold is deliberately left to the clinic: the delivered calibration is a single near-balanced patient (P11), and a larger dedicated calibration set would tighten it.

5.3. Future Research Directions

Several directions follow naturally. The full annotated corpus can now train a final model under the P5-A protocol at expanded scale, and the cls weighting and freeze depth, fixed early when each fold trained on roughly 200 images, are worth re-ablating now that folds train on over 1,000 positives. A dedicated, class-balanced held-out calibration set, larger than the single P11 patient, would give the hospital a cleaner out-of-sample threshold than a single patient can. The clearest validity step is external validation on aspirates from a second institution, which would test the SM-enriched training distribution against genuinely independent staining and scanning. On the deployment side, integration into Cinderella, the cell-classification software already in use at USZ, becomes worthwhile once the institutional security and change-management timeline permits, turning the standalone tool into part of the routine scan-to-inference-to-evaluation pipeline.

6. Declaration on AI Assistance

The author declares that the empirical research, data analysis, experimental design, and quantitative results presented in this thesis were generated independently and strictly remain the intellectual property of the author.

Artificial intelligence (AI) assistance was utilized solely for the purpose of language refinement, stylistic editing, optimization of academic reading flow and coding assistance. Specifically, the AI models (Gemini 2.5 Pro, Claude Sonnet 4.5, and Claude Sonnet 4.6) were employed to translate text, correct grammar, and restructure sentences into the required third-person past tense and formal academic style.

No core scientific content, model selection decisions, hyperparameter choices, or interpretation of results were delegated to or generated by the AI model.

7. References

- [1] T. Ghete *et al.*, “Models for the marrow: A comprehensive review of AI-based cell classification methods and malignancy detection in bone marrow aspirate smears,” *HemaSphere*, vol. 8, no. 12, p. e70048, 2024, doi: <https://doi.org/10.1002/hem3.70048>.
- [2] H. Sazak and M. Kotan, “Automated blood cell detection and classification in microscopic images using YOLOv11 and optimized weights,” *Diagnostics*, vol. 15, no. 1, 2025, doi: [10.3390/diagnostics15010022](https://doi.org/10.3390/diagnostics15010022).
- [3] J. Diaz, A. Deb, A. Lyubimova, L. Santos-Carrion, C. Romagnoli, and J. Barrientos, “Machine learning-based detection model for leukemia cells in peripheral blood smears using YOLOv11-large,” *Blood*, vol. 146, p. 6123, 2025, doi: <https://doi.org/10.1182/blood-2025-6123>.
- [4] G. Jocher, J. Qiu, and A. Chaurasia, *Ultralytics YOLO*. (Jan. 2023). Available: <https://github.com/ultralytics/ultralytics>
- [5] F. Vollmer, “MastCellDetector: Desktop application for mast cell detection and classification.” 2026. Available: https://github.com/kekslol24/desktop_app
- [6] F. Vollmer, “Automated detection and morphological classification of mast cells in bone marrow aspirates using deep learning.” GitHub, 2026. Available: <https://github.com/kekslol24/BA>
- [7] P. Valent *et al.*, “Updated diagnostic criteria and classification of mast cell disorders: A consensus proposal,” *HemaSphere*, vol. 5, no. 11, p. e646, Oct. 2021, doi: [10.1097/HS9.0000000000000646](https://doi.org/10.1097/HS9.0000000000000646).
- [9] Gemini 2.5 Pro. (n.d.). Google Gemini. Retrieved October 1, 2025, from <https://gemini.google.com/app>
- [10] DeepL Translate: The world’s most accurate translator. (n.d.). <https://www.deepl.com/>

8. Lists

List of Figures

Figure 1	Analysis tab of the Gradio prototype showing uploaded image tiles with YOLO detection overlays.	18
Figure 2	Research tab showing ground-truth annotations alongside model predictions for direct visual comparison.	18
Figure 3	Research tab metrics output: precision-recall curves and performance plots generated from the ground-truth evaluation run.	19
Figure 4	Statistics dashboard of the Gradio prototype showing per-class detection counts and the Atypisch fraction relative to the WHO 25% diagnostic threshold.	19
Figure 5	Gallery view of the desktop application showing all image tiles for a patient folder with sidebar controls for inference configuration.	20
Figure 6	Gallery filtered to tiles with detections, allowing rapid navigation to images where the model found mast cells.	21
Figure 7	Annotation editor with interactive bounding boxes; boxes can be added, moved, resized, or deleted directly in the application without external tooling.	22
Figure 8	Statistics tab aggregating per-image detections into a patient-level summary with the Atypisch fraction displayed against the WHO 25% diagnostic threshold.	23
Figure 9	Training convergence for P5-B Fold 3 (G3 holdout: P9, P12, P53). Panels from top left: box loss, classification loss, and distribution focal loss for training and validation splits; precision, recall, mAP50, and mAP50-95 across epochs.	33
Figure 10	Normalized confusion matrix for P5-B Fold 3 (G3 holdout: P9, P12, P53). Rows are ground-truth classes; columns are predicted classes. Values are row-normalized fractions. The main diagonal entries are 0.96 (Atypisch) and 0.88 (Normal); cross-class error is 0.01; background-to-Atypisch is 0.91.	34
Figure 11	P5-B per-group recall by class. Bars show R_Atypisch (blue) and R_Normal (orange) for each held-out patient group. Dashed horizontal lines mark the cross-fold means (R_Atypisch = 0.885, R_Normal = 0.821).	35
Figure 12	Performance comparison across P5-A (35-fold LOPO, test), P5-B (5-group CV, test), and P6-A (final model, val). Error bars show ± 1 std across folds, clipped at 1.0. P6-A is a single training run with no held-out test set.	35
Figure 13	Metric comparison between P5-A (35-fold LOPO) and P5-B (grouped 5-fold CV). Error bars represent ± 1 standard deviation across folds. R_Atyp and R_Norm for P5-	

A are computed over the 20 and 29 folds respectively where the class was present in the test set.	36
Figure 14 Training curves for the P6-A final deployment model. Panels show box loss, classification loss, and DFL loss for training and validation splits; precision, recall, mAP50, and mAP50-95 across epochs.	36
Figure 15 Normalized confusion matrix for P6-A (best.pt, evaluated on the val split). Rows are ground-truth classes; columns are predicted classes.	37
Figure 16 Confidence distribution of the final deployment model (P6-A) across the complete P11 tile set (~12'000 tiles; 17 Atypisch and 20 Normal annotated).	38

List of Tables

Table 1 Full corpus summary across P1-P53.	10
Table 2 Milestone run progression summary across all experimental phases. All metrics shown as mean $\pm$ std across folds except P6-A (single training run, val split only; no held-out test set). Base model is yolo11n throughout except P3-E (yolo11l, 43 M parameters). CV-codes: S = stratified k-fold, L = leave-one-patient-out (LOPO), CV = grouped cross-validation; prefix = fold/group count. *P3-A per-class recall values are cross-fold means from the tinkering matrix reference configuration (degrees=0, dfl=1.5, freeze=10); folds where the respective class was absent from the held-out test set are excluded from each per-class average.	28
Table 3 P5-A per-fold test results (35-fold LOPO post-hoc evaluation, full P1–P53 annotated corpus). GT Atyp / GT Norm = ground-truth cell instance counts in the held-out test split. “—” indicates the class was absent from that test split and is excluded from the per-class mean. R_Atyp mean over n=20 folds; R_Norm mean over n=29 folds.	30
Table 4 P5-B grouped 5-fold CV split composition and test results (35 annotated patients, P1–P53 corpus). Train = annotated (positive) images; each fold additionally receives 3,773-4,595 negative tiles. Atyp/Norm = ground-truth label counts in the test split. All five groups have both classes present.	32

9. Attachments